

While collaborating on a research project with colleagues from the University of Antwerp, we became interested in understanding how frequently specific amino-acid motifs appear in the human proteome. This question arose because we were studying a neuropathy-causing mutation in the protein HSP27, which is a small heat-shock protein that functions as a molecular chaperone. The particular mutation in HSP27 changes the native proline residue at position 182 to leucine (P182L), which disrupts a conserved motif known as the "[I/V]x[I/V]" motif wherein the first and third residues are either isoleucine (I) or valine (V) while the central residue can be anything (x). In mammalian small heat-shock proteins, the central residue is typically proline.

The [I/V]x[I/V] motif is an example of a short linear motif (SLiM), which are motifs of typically 3-8 residues that can play important biological roles, such as mediating protein-protein interactions or promoting degradation. Notably, SLiMs generally appear in intrinsically disordered regions (IDRs) of proteins, thus ensuring accessibility and promoting promiscuous interactions with many other proteins. If you are interested in searching for SLiMs in your protein of interest, here is a [useful website called SlimSearch](#), which is based on [this paper from the Davey Lab](#).

Anyway, during our hunt for [I/V]x[I/V] SLiMs in the human proteome, I wrote some Python scripts that I used to perform proteome-wide calculations on UniProt proteomes or FASTA files. I've written a short Python tutorial that outlines the following steps, with example screenshots and code showing how to:

1. Download and format UniProt proteome files
2. Compute the mean and median protein length
3. Plot histograms of protein length distributions

I wrote this in Python3.7.4 and the code requires a few packages (Numpy, Pandas, and BioPython) that you might have to install ahead of time.

First, open up a new text file and import these packages

```
1 import os, sys
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import pandas as pd
5 from Bio import SeqIO
6 import itertools
```

Now, let's create a class called **Proteome** where we will perform our large-scale counting exercises. By default, we must begin the class with the init function. In case you are unfamiliar with classes, take a look at this [helpful explanation](#).

```
1 class Proteome(object):
2     """ Some useful calculations on proteomes """
3
4     def __init__(self):
5         self.type = None
```

Save this file as "run.py" in the directory of your choice.

In order to start counting motifs within a proteome, we first need a proteome! The UniProt website maintains a large repository of proteomes, so download the **Reviewed** human proteome from [this link](#). Make sure to download the files in **FASTA** format! UniProt maintains a very large database of alternatively spliced isoforms of each gene; for the purposes of this example, we are only interested in the "canonical" isoform of each gene, i.e. one gene, one amino acid sequence. So, make sure to click the "Reviewed" version of the proteome.

UniProt

Proteomes

Advanced Search

BLAST Align Retrieve/ID mapping Peptide search Help Contact

Proteomes - Homo sapiens (Human)

Overview

None

☒ Overview

☒ Components

☒ Publications

Map to

UniProtKB (74,811)

☒ Reviewed (20,350)

☐ Unreviewed (54,461)

Swiss-Prot

TrEMBL

Status	Reference proteome
Proteins	74,811
Gene count	20,595 - Download one protein sequence per gene (FASTA)
Proteome ID	UP000005640
Taxonomy	9606 - Homo sapiens
Last modified	December 1, 2019
Genome assembly	GCA_000001405.27 from Ensembl full
Busco	C:99.5%[S:38.2%,D:61.3%],F:0.2%,M:0.3%,n:6192
Completeness	Outlier

Homo sapiens (*Homo sapiens sapiens*) or modern humans are the only living species of the evolutionary branch of great apes known as hominids. Divergence of early humans from chimpanzees and gorillas is estimated to have occurred between 4 and 8 million years ago. The genus *Homo* (*Homo habilis*) appeared in Africa around 2.3 million years ago and shows the first signs of stone tool usage. The exact lineage of *Homo* species ie: *H. habilis*/*H. ergaster* to *H. erectus* to *H. rhodesiensis*/*H. heidelbergensis* to *H. sapiens* is still hotly disputed. However, continuing evolution and in particular larger brain size and complexity culminates in *Homo sapiens*. The first anatomically modern humans appear in the fossil record around 200,000 years ago. Modern humans migrated across the globe essentially as hunter-gatherers until around 12,000 years ago when the practice of agriculture and animal domestication enabled large populations of civilizations. Overall life expectancy in Europe is 81 years.

[https://www.uniprot.org/uniprot/?query=proteome:UP000005640 reviewed:...](https://www.uniprot.org/uniprot/?query=proteome:UP000005640 reviewed:)

Figure 1: [Navigate to the human proteome on UniProt and click **Reviewed**.](#)

After clicking on **Reviewed**, the website will take you to the page shown below. However, before we download the proteome, we first must specify a few things about the file format. Click on the **Columns** button to specify the data that we'll include in the file.

BLAST Align Retrieve/ID mapping Peptide search Help Contact

UniProtKB results

UniProtKB consists of two sections:

- Reviewed (Swiss-Prot) - Manually annotated**
Records with information extracted from literature and curator-evaluated computational analysis.
- Unreviewed (TrEMBL) - Computationally analyzed**
Records that await full manual annotation.

The UniProt Knowledgebase (UniProtKB) is the central hub for the collection of functional information on proteins, with accurate, consistent and rich annotation. In addition to capturing the core data mandatory for each UniProtKB entry (mainly, the amino acid sequence, protein name or description, taxonomic data and citation information), as much annotation information as possible is added.

Help UniProtKB help video Other tutorials and videos Downloads

Filter by: Reviewed (20,350) Swiss-Prot Popular organisms Human (20,350) Proteomes UP000005640 (20,350) View by

BLAST Align Download Add to basket Columns

1 to 25 of 20,350 Show 25

Entry	Entry name	Protein names	Gene names	Organism	Length
Q9NQ39	RS10L_HUMAN	Putative 40S ribosomal protein S10...	RPS10P5 RPS10L	Homo sapiens (Human)	176
Q99470	SDF2_HUMAN	Stromal cell-derived factor 2	SDF2	Homo sapiens (Human)	211
Q8N8B7	TEANC_HUMAN	Transcription elongation factor A N...	TCEANC	Homo sapiens (Human)	351
Q658L1	SAXO2_HUMAN	Stabilizer of axonemal microtubules...	SAXO2 FAM154B	Homo sapiens (Human)	398
O75494	SRS10_HUMAN	Serine/arginine-rich splicing facto...	SRSF10 FUSIP1, FUSIP2, SFRS13A, TASR	Homo sapiens (Human)	262

Figure 2: [Specify which parameters to include in the download file by clicking on *Columns*.](#)

On the **Column** page, you must click the box next to **Sequence** in order to indicate that you want each protein's sequence included in the downloaded file.

Search:

Names & Taxonomy

- ☒ Entry name
- ☒ Gene names
- ☐ Gene names (ordered locus)
- ☐ Gene names (ORF)
- ☐ Gene names (primary)
- ☐ Gene names (synonym)
- ☒ Organism
- ☐ Organism ID
- ☒ Protein names
- ☐ Proteomes
- ☐ Taxonomic lineage
- ☐ Virus hosts

Sequences

- ☐ Alternative products (isoforms)
- ☐ Alternative sequence
- ☐ Erroneous gene model prediction
- ☐ Fragment
- ☐ Gene encoded by
- ☒ Length
- ☐ Mass
- ☐ Mass spectrometry
- ☐ Natural variant
- ☐ Non-adjacent residues
- ☐ Non-standard residue
- ☐ Non-terminal residue
- ☐ Polymorphism
- ☐ RNA editing
- ☒ Sequence
- ☐ Sequence caution
- ☐ Sequence conflict
- ☐ Sequence uncertainty
- ☐ Sequence version

Function

- ☐ Absorption
- ☐ Active site
- ☐ Activity regulation
- ☐ Binding site
- ☐ Calcium binding
- ☐ Catalytic activity
- ☐ Cofactor
- ☐ DNA binding
- ☐ EC number
- ☐ Function [CC]ⁱ
- ☐ Kinetics
- ☐ Metal binding
- ☐ Nucleotide binding
- ☐ Pathway
- ☐ pH dependence
- ☐ Redox potential
- ☐ Rhea Ids
- ☐ Site
- ☐ Temperature dependence

Miscellaneous

- ☐ Annotation
- ☐ Caution
- ☐ Features
- ☐ Keyword ID
- ☐ Keywords
- ☐ Matched text
- ☐ Miscellaneous [CC]ⁱ
- ☐ Protein existence
- ☒ Tools
- ☐ UniParc

Interaction Expression Gene Ontology (GO) Chemical entities (ChEBI)

Figure 3: [Click the box next to *Sequence*](#)

Then scroll up and click **Save** in the upper-right corner.

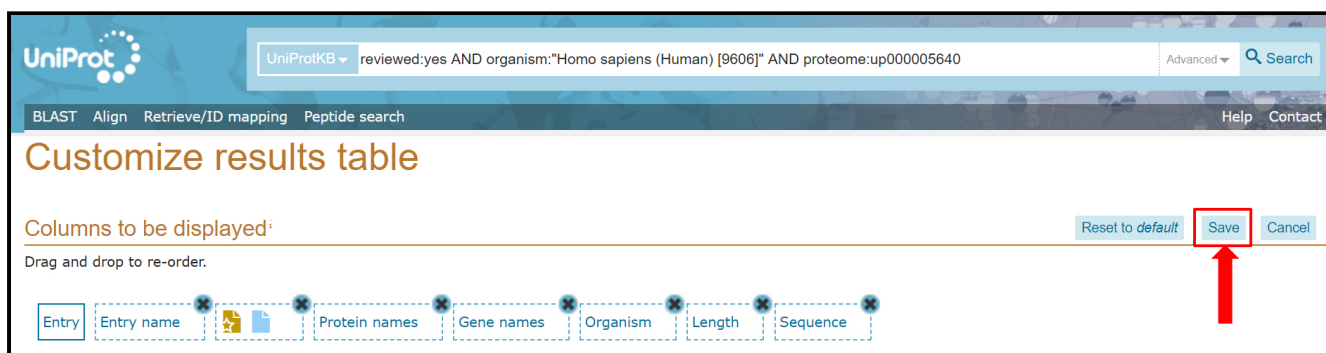


Figure 4: [Click Save](#)

This will take you back to the UniProtKB Results screen from above. Click on **Download** and then, under **Format**, choose **Tab-separated** from the drop-down menu. Also, make sure to click the button **Uncompressed**. Now, you can click **Go** to download the file containing 20,350 FASTA files. I saved this file as "uniprot_human.tab" and it's roughly 14.5 MB in size.

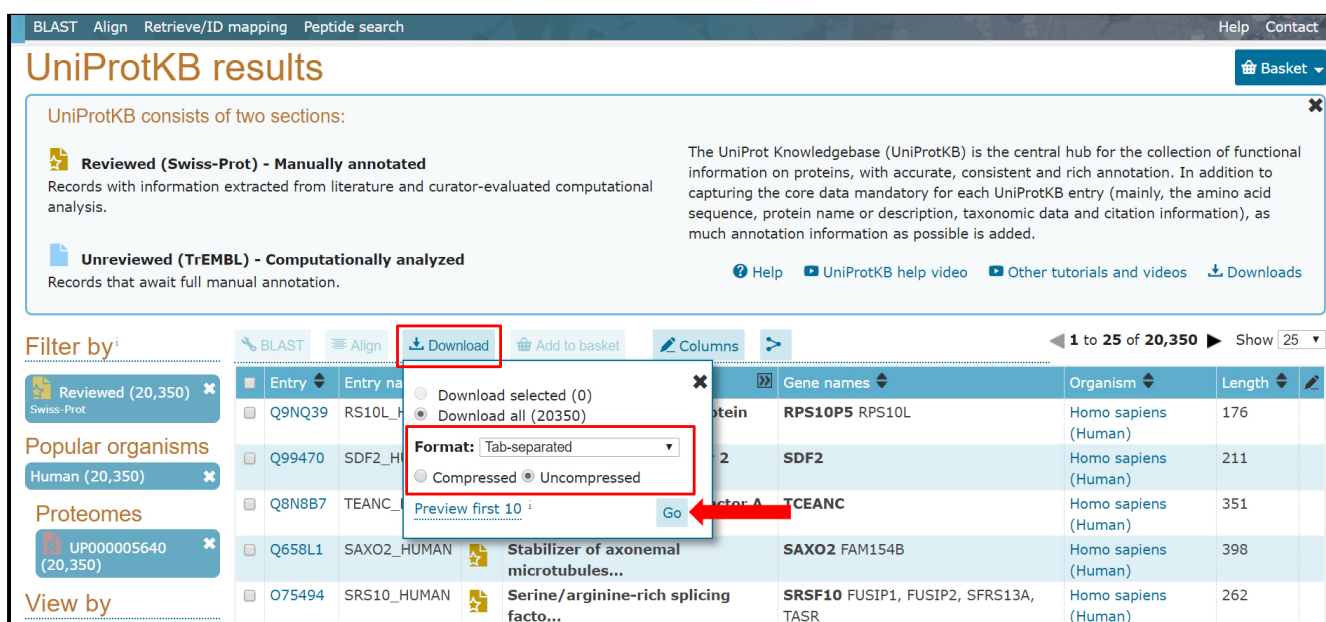


Figure 5: [Download the tab-separated file](#)

Before we move ahead, let's open up the file and see what it contains. It should come as a [long-name-here].tab file, so rename it something convenient and store it in the same directory as your new .py file. Then, open up the .tab file in your favorite text editor. Below is a screenshot from Notepad++, and you can see that I've highlighted the first row that contains the column headings: **Entry, Entry name, Status, Protein names, Gene names, Organism, Length, Sequence**. At the bottom of Notepad++, it shows us that we have 20,352 lines in this file. The first line contains the column headers and the very last line is empty, so there are **20,350 proteins** here. This is a good sanity check: the number of proteins in the file matches the value listed on the UniProt website.

Entry	Entry name	Status	Protein names	Gene names	Organism	Length	Sequence
Q9NQ39	RS10L_HUMAN	reviewed	Putative 40S ribosomal protein S10-like RPS10P5 RPS10L		Homo sapiens (Human)	176	MLMPKKNRIAIHELHFKEGVMVAKKDVMHPKHPE
Q99470	SDF2_HUMAN	reviewed	Stromal cell-derived factor 2 (SDF-2)	SDF2	Homo sapiens (Human)	211	MAVVEPLLLGGGLNSAVGASSLGVVTCGSSVVKLLNTRHNVRLH
Q8NB7	TEANC_HUMAN	reviewed	Transcription elongation factor A N-terminal and central domain-containing protein (TFIIIS central domain-containing pr				
Q658L1	SAXO2_HUMAN	reviewed	Stabilizer of axonemal microtubules 2	SAXO2 FAM154B	Homo sapiens (Human)	398	MGAKSMRSWCLQCICSGSDYCPYEIVKQPRHVP
Q75494	SRS10_HUMAN	reviewed	Serine/arginine-rich splicing factor 10 (40 kDa SR-repressor protein) (SRrp40) (FUS-interacting serine-arginine-rich p				
Q9NQ55	SSF1_HUMAN	reviewed	Suppressor of SWI4 1 homolog (Ssf-1) (Brix domain-containing protein 3) (Peter Pan homolog) PPAN BXDC3 SSF1 Homo sapie				
Q8IW03	SIAH3_HUMAN	reviewed	Seven in absentia homolog 3 (Siah-3)	SIAH3	Homo sapiens (Human)	269	MLFFTQCQFAGVLDLIHLRFQHYKAKRVFSAAGQLVCVNPPTH
Q5HYW3	RTL5_HUMAN	reviewed	Retrotransposon Gag-like protein 5 (Retrotransposon gag domain-containing protein 4)	RTL5 KIAA2001 RGAG4	Homo sapie		
Q96L92	SNX27_HUMAN	reviewed	Sorting nexin-27	SNX27 KIAA0488 My014	Homo sapiens (Human)	541	MADEDGEGIHFSAPHNRNGGGGGGGSLHCAGNGGGGGGGPRVVR
Q99598	TSNAX_HUMAN	reviewed	Translin-associated protein X (Translin-associated factor X)	TSNAX TRAX	Homo sapiens (Human)	290	MSNKEGGGFRKRR
Q9H489	TSY26_HUMAN	reviewed	Putative testis-specific Y-encoded-like protein 3 (TSPY-like protein 3) (Testis-specific Y-encoded protein 26 pseudoge				
Q9765	VEGFB_HUMAN	reviewed	Vascular endothelial growth factor B (VEGF-B) (VEGF-related factor) (VRF)	VEGFB VRF	Homo sapiens (Human)	207	MS
Q96EX3	WDR34_HUMAN	reviewed	WD repeat-containing protein 34	WDR34	Homo sapiens (Human)	536	MATRAQPGPLSQAGSAGVAALATVGASGPGRPGPLQDETGLVASVPS
Q9NYU1	UGGG2_HUMAN	reviewed	UDP-glucose:glycoprotein glucosyltransferase 2 (UGT2) (hUGT2) (EC 2.4.1.-) (UDP--Glc:glycoprotein glucosyltransferase				
Q8NDW8	TT21A_HUMAN	reviewed	Tetratricopeptide repeat protein 21A (TPR repeat protein 21A) (Stress-inducible protein 2)	TT21A STI2	Homo sapiens (
Q5JXB2	UBE2NL_HUMAN	reviewed	Putative ubiquitin-conjugating enzyme E2 N-like (Epididymis tissue protein Li 174)	UBE2NL	Homo sapiens (Human)	15	
Q9NR4	SDHF3_HUMAN	reviewed	Succinate dehydrogenase assembly factor 3, mitochondrial (SDH assembly factor 3) (SDHAF3)	SDHAF3	SDHAF3 ACN9 DC11	Homo s	
Q94915	TSC2_HUMAN	reviewed	Tuberlin (Tuberous sclerosis 2 protein)	TSC2 TSC4	Homo sapiens (Human)	1807	MAKPTSKDGLKEFKILLGLSTFRPNPRAEGK
Q96EV3	TSPAN17_HUMAN	reviewed	Tetraspanin-17 (Tspan-17) (F-box only protein 23) (Tetraspan protein SB134) (Transmembrane 4 superfamily member 17) TS				
Q95859	TSPAN12_HUMAN	reviewed	Tetraspanin-12 (Tspan-12) (Tetraspan NET-2) (Transmembrane 4 superfamily member 12) TSPAN12 NET2 TM4SF12 UNQ774/PRO156				
Q05066	SRX_HUMAN	reviewed	Sex-determining region Y protein (Testis-determining factor)	SRX TDF	Homo sapiens (Human)	204	MQSYASAMLSVFNDDYS
Q9Y2R5	RT17_HUMAN	reviewed	28S ribosomal protein S17, mitochondrial (MRP-S17) (S17mt) (Mitochondrial small ribosomal subunit protein uS17m) MR				
Q969E3	UCN3_HUMAN	reviewed	Urocortin-3 (Stresscopin) (Urocortin III) (Ucn III) UCN3 SPC		Homo sapiens (Human)	161	MLMPFVHLLLLLLLGGPRTGLPHKF
Q5T1N1	AKNAD1_HUMAN	reviewed	Protein AKNAD1 AKNAD1 Clorf62		Homo sapiens (Human)	836	MDEADFSEHTTYKQEDLPYDGLDSQIKIGNDYFTSKKDGLEVLNQIIFIADDPQEKW
P17213	BPI_HUMAN	reviewed	Bactericidal permeability-increasing protein (BPI) (CAP 57) BPI		Homo sapiens (Human)	487	MRENMARGPCNAPFWASIMLVVAIGT
Q92772	CDKL2_HUMAN	reviewed	Cyclin-dependent kinase-like 2 (EC 2.7.11.22) (Protein kinase p56 KKIAMRE) (Serine/threonine-protein kinase KKIAMRE)				
Q9P0S2	COX16_HUMAN	reviewed	Cytochrome c oxidase assembly protein COX16 homolog, mitochondrial (hCOX16)	COX16 C14orf112 HSPC203 PTD019	Homo sapie		
Q92685	ALG3_HUMAN	reviewed	Dol-P-Man:Man(5)GlcNAc(2)-PP-Dol alpha-1,3-mannosyltransferase (EC 2.4.1.258) (Asparagine-linked glycosylation protein				
P15088	CBPA3_HUMAN	reviewed	Mast cell carboxypeptidase A (MC-CPA) (EC 3.4.17.1) (Carboxypeptidase A3)	CPA3	Homo sapiens (Human)	417	MRLLIF
Q96HB5	CC120_HUMAN	reviewed	Coiled-coil domain-containing protein 120	CCDC120 JM11	Homo sapiens (Human)	630	MEVKQLISSPTFNAPALFGEAAPQVKSE
Q95415	BRI3_HUMAN	reviewed	Brain protein I3 (pRGR2)	BRI3	Homo sapiens (Human)	125	MDHKPLLQERPPAYNLEAGQGYACGPHGYGAIPAAPPPPPYPLVTGIPTHHP
H3BNL1	CC084_HUMAN	reviewed	Uncharacterized protein C3orf84 C3orf84		Homo sapiens (Human)	204	MQSALVGSWHNNFGYGHYRSQFKSESAREYHLAAKQPPAVFLGRCQEPV
P16152	CBRI1_HUMAN	reviewed	Carbonyl reductase [NADPH] 1 (EC 1.1.1.184) (15-hydroxyprostaglandin dehydrogenase [NADP(+)]) (EC 1.1.1.197) (NADPH-de				
Q9H8M2	BRD9_HUMAN	reviewed	Bromodomain-containing protein 9 (Rhabdomyosarcoma antigen MU-RMS-40.8) BRD9 UNQ3040/PRO9856		Homo sapiens (Human)		
Q177F1	CCB42_HUMAN	reviewed	Putative chemokine-related protein B42		Homo sapiens (Human)	81	MPLSDWCCGICEEAPLGRAYTQTWMTGCGPHGVTALGQQLKDCDL

Figure 6: Open up the file and peruse it

Now that we know what our proteome file looks like, we can write a Python function to open it up and store it in a Pandas dataframe.

```

1 def load_uniprot(self, uniprot):
2     """ Reads a UniProt file; returns a pandas dataframe
3
4     Parameters
5     -----
6     uniprot : file downloaded from UniProt containing a proteome of interest
7
8     Returns
9     -----
10    proteome_dataframe, a Pandas dataframe that contains the FASTA IDs, sequences,
    and length of sequences """
11
12    self.uniprot_proteome = pd.read_csv(uniprot, sep='\t', names=['Entry', 'Entry
    name', 'Status', 'Protein names', 'Gene names', 'Organism', 'Length', 'Sequence'])
13    self.uniprot_proteome['Length'][1:] = self.uniprot_proteome['Length']
    [1:].astype(int) # convert the length values to integers
14
15    return self.uniprot_proteome

```

We've now loaded the proteome into a Pandas dataframe that contains 3 columns:

1. **ID** = the UniProt ID
2. **Sequence** = the protein's sequence
3. **Length** = the protein's length.

To read the fasta file, we've made use of the [SeqIO.index function in BioPython](#), which converts a FASTA file into a dictionary with keys and values respectively corresponding to IDs and sequences. With **SeqIO.index()**, one has a fast way to load very large FASTA files into Python without having to do a loop or loading of all lines into memory at once (e.g. with the Seq.list() or Seq.to_dict() functions).

With our proteome in an accessible data structure, we can begin to calculate some basic properties. For instance, we might want to know the total proteome length, in case we'd like to compute percentages.

```
1  def proteome_length(self, query_proteome):
2      """ Reads a Pandas dataframe containing a proteome of interest; returns the
   total number of residues
3
4      Parameters
5      -----
6      query_proteome : Pandas dataframe, contains the proteome of interest
7
8      Returns
9      -----
10     total_length : int, total number of residues in proteome_dataframe """
11
12     total_length = query_proteome['Length'][1:].sum(axis=0, skipna=True)
13
14     return total_length
```

This function takes a proteome in a Pandas dataframe as input and **returns the total number of residues in that proteome**. To quickly count up the number of residues in the proteome, we can sum the **Length** column that we've already appended to our dataframe in the **load_fasta** function above. The argument "axis=0" is required to tell Pandas to sum down the column and not across the row.

So, now that we have the total number of residues in the proteome, let's write a function to search the proteome for a specific motif. We want this function to output two values:

1. the total number of times that we counted the queried motif. Let's call this value **sum**.
2. the observed number of motifs divided by the total number of peptides of size n in the proteome, or the ***fraction***.

Because we are counting up motifs in a "sliding" manner, *i.e.* by starting at residue 1 and working our way toward the C-terminus, then we should also calculate how many peptides of size n exist in a similar manner. For example, if our queried motif is "QVERY" then our denominator, x , becomes:

$$x = (N-n)+1$$

where N is the length of the proteome and n is the length of the queried motif. (I encourage you to calculate this on your own using, e.g. $N = 10$ and $n = 2$. You'll see that there exist 9 di-peptides when sliding from residue 1 to residue 10, *i.e.* 1-2, 2-3, 3-4, ..., 9-10).

```

1  def find_motifs(self, data_frame, query):
2      """ Searches a Pandas dataframe containing a proteome for a specific amino
3      acid/motif; returns the number of instances and fraction
4
5      Parameters
6      -----
7      data_frame : Pandas dataframe, contains the proteome of interest
8      query : str, the amino acid or motif to be searched
9
10     Returns
11     -----
12     sum : int, total number of hits for the queried residue/motif
13     fraction: float, sum divided by the total number of residues or motifs of the
14     same size """
15
16     # Extract sequences from the pandas dataframe
17     seqs = data_frame['Sequence']
18
19     # Calculate the total number of residues in the queried proteome
20     total_residues = self.proteome_length(data_frame)
21
22     # Initialize an array
23     hits_array = []
24
25     # Loop over each sequence and count the number of times the queried motif
26     shows up
27     for protein in seqs:
28         hits_array.append(protein.count(query))
29
30     # total number of hits for the queried residue/motif, check if any hits are
31     found
32     sum = np.sum(np.asarray(hits_array))
33     if sum < 1:
34         fraction = 0
35     else:
36         fraction = sum/(total_residues - len(query)+1) # total number of hits
37         divided by the total number of residues or motifs of the same size
38
39     return sum, fraction

```

Let's test out the code. We can quickly instantiate the class, load the FASTA file, and calculate the total length of the proteome.

```
1 # Instantiate the class
2 human = Proteome()
3
4 # Load the FASTA file containing the human proteome
5 uniprot = human.load_uniprot('uniprot_human.tab') # replace with the name of your .tab
  file
6
7 # What is the length of the proteome?
8 print('Proteome length = %.0f AA' % human.proteome_length(uniprot))
```

```
Proteome length = 11354232 AA
```

For the human proteome, we see that there are $N = 11,354,232$ residues. So, for our example above with the motif "QVERY", there are $x = (11,354,232 - 5) + 1 = 11,354,228$ penta-peptides in the proteome.

With our new function **find_motifs**, we can also count how many particular motifs exist in the proteome.

```
1 # Find how many Q, QV, QVE, QVER, QVERI, QVERIE, QVERIES
2 motif = 'QVERIES'
3 for i in range(0, len(motif)+1):
4     print('%s = %.0f' % (motif[0:i], human.find_motifs(uniprot, motif[0:i])[0]))
```

```
> Q = 541533
> QV = 31622
> QVE = 2108
> QVER = 116
> QVERI = 6
> QVERIE = 2
> QVERIES = 0
```

By iterating over the queried motif "QVERIES", we see that the number of found motifs dramatically decreases with increasing length of the peptide. In fact, in the entire proteome comprising 11,354,226 hepta-peptides, **there is not a single QVERIES motif to be found**. We'll return to this point a little later.

We now know that there are about **11.4 million amino acids** in the proteome. Indeed, it would take some time to count motifs by hand...! We can do some basic calculations with these numbers, and we find that **the average length of a human protein is 558 residues**, or *ca* 61 kDa (using an estimate of *ca* 110 Da per amino acid).


We can also compute this with numpy, where we obtain the same average length of 558 residues. The median length, however, is only 415 residues, and this reflects a skewed distribution.

```
1 # What are the mean & median protein lengths in the proteome?
2 mean = np.mean(uniprot['Length'][1:])
3 median = np.median(uniprot['Length'][1:])
4 print('Mean protein length = %.0f AA' % mean)
5 print('Median protein length = %.0f AA' % median)
```

```
> Mean protein length = 558 AA
> Median protein length = 415 AA
```

We can plot a histogram of protein lengths to visualize the distribution over the proteome. I limited the x axis because the vast majority of proteins are fewer than 3000 residues. Indeed, we can see that our distribution of protein sizes (below) are skewed to the right with a tail.

```
1 # Plot a histogram of sequence length
2 plt.hist(uniprot['Length'][1:], bins=1000, color='k', edgecolor='w', linewidth=0.1)
3 plt.title('Protein size distribution in the human proteome')
4 plt.ylabel('Count')
5 plt.xlabel('Number of amino acids')
6 plt.xlim(0,3000)
7 plt.tight_layout()
8 plt.show()
```



Histogram of protein lengths in the human proteome

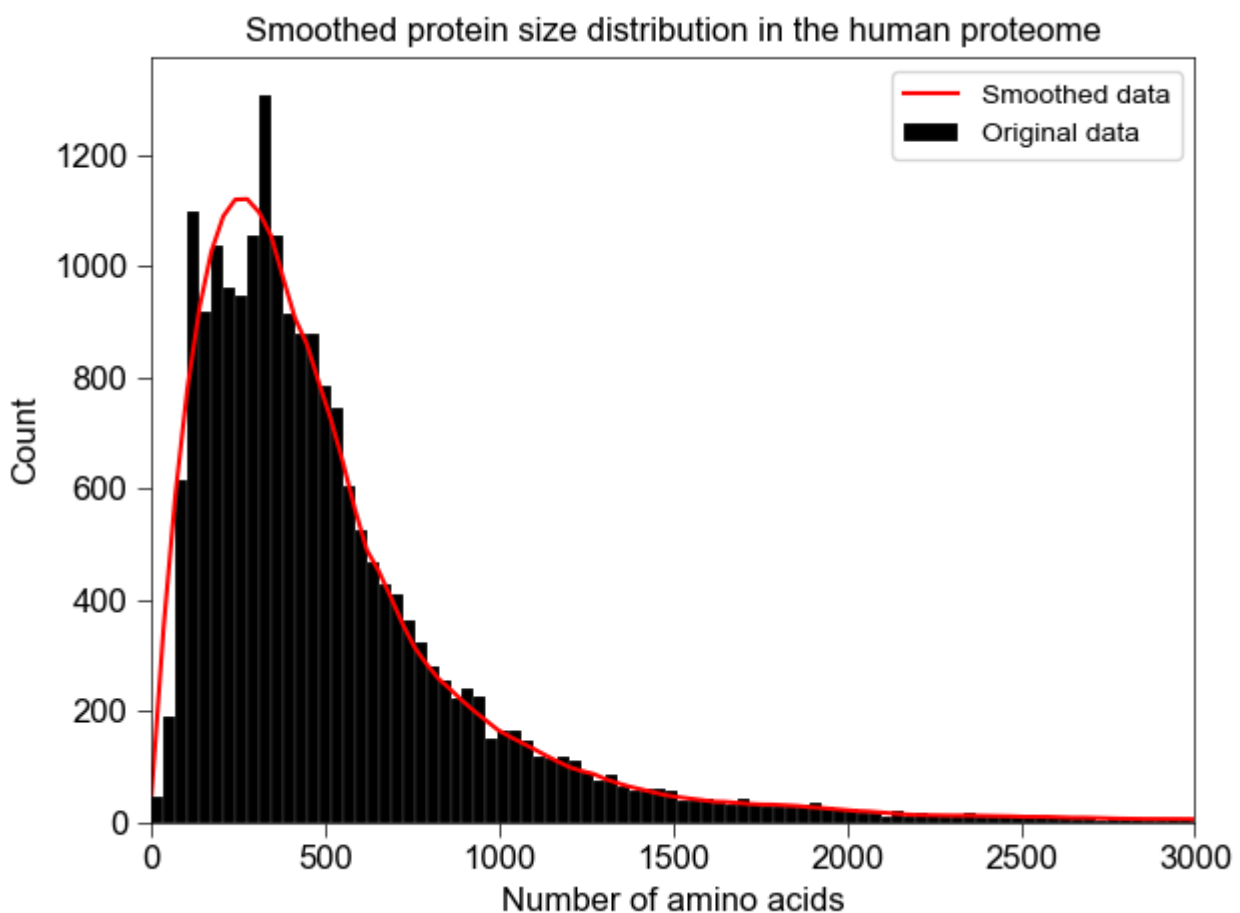
We can smooth out our distribution by applying a filter known as the [Savitzky-Golay filter](#).

```
1  from scipy.signal import savgol_filter
2
3  def filter(data_to_filter, window=19, poly=3):
4      """ Applies the Savitzky-Golay filter to a 1D array
5
6      Parameters
7      -----
8      data_to_filter : dataframe or array, contains the data to be smoothed
9      window : int, the length of the filter window
10     poly : int, the order of the polynomial (must be < window)
11
12     Returns
13     -----
14     filtered : array, the smoothed data """
15     data = np.asarray(data_to_filter).astype(np.float64)
16     filtered = savgol_filter(data, window, poly)
17
18     return filtered
19
20 # Bin the data
```

```

21 data, bins = np.histogram(np.asarray(uniprot['Length'][1:]), bins=1000, density=False)
22
23 # Smooth the binned data
24 filtered_data = filter(data, window=19, poly=3)
25
26 # Plot the histogram & the smoothed data
27 plt.hist(uniprot['Length'][1:], bins=1000, color='k', edgecolor='w', linewidth=0.1,
28          label='Original data')
29 plt.plot(bins[:-1], filtered_data, 'r', label='Smoothed data')
30 plt.title('Smoothed protein size distribution in the human proteome')
31 plt.ylabel('Count')
32 plt.xlabel('Number of amino acids')
33 plt.xlim(0,3000)
34 plt.ylim([0,None])
35 plt.legend(loc='upper right')
36 plt.tight_layout()
37 plt.show()

```




Smoothing the histogram with a Savitzky-Golay filter

We can also recover the histogram-style plot of our smoothed data by making use of the **weights** variable in the `plt.hist` function:

```

1 # Plot our histogram as above
2 plt.hist(uniprot['Length'][1:], bins=1000, density=False, color='k', alpha=0.5,
3          edgcolor='w', linewidth=0.1, label='Original data')
4
5 # Plot the smoothed data as a histogram by setting weights = filtered_data
6 plt.hist(bins[:-1], bins=1000, weights = filtered_data, density=False, color='r',
7          alpha=0.5, edgcolor='w', linewidth=0.1, label='Smoothed data')
8
9 # These are the exact same plotting parameters as above
10 plt.title('Smoothed protein size distribution in the human proteome')
11 plt.ylabel('Count')
12 plt.xlabel('Number of amino acids')
13 plt.xlim(0,3000)
14 plt.ylim([0,None])
15 plt.legend(loc='upper right')
16 plt.tight_layout()
17 plt.show()

```



Smoothing the histogram with a Savitzky-Golay filter and recovering the binned values

Now that we've plotted our smoothed distribution, we can compare this to other organisms. Do humans, on average, have shorter or longer proteins than some model organisms like *E. coli* or *S. cerevisiae*? To test this, we can simply download the respective proteomes from UniProt using the steps outlined above. You can click these links to reach [the *E. coli* proteome](#) or [the *S. cerevisiae* proteome](#) that I used for the calculations below.

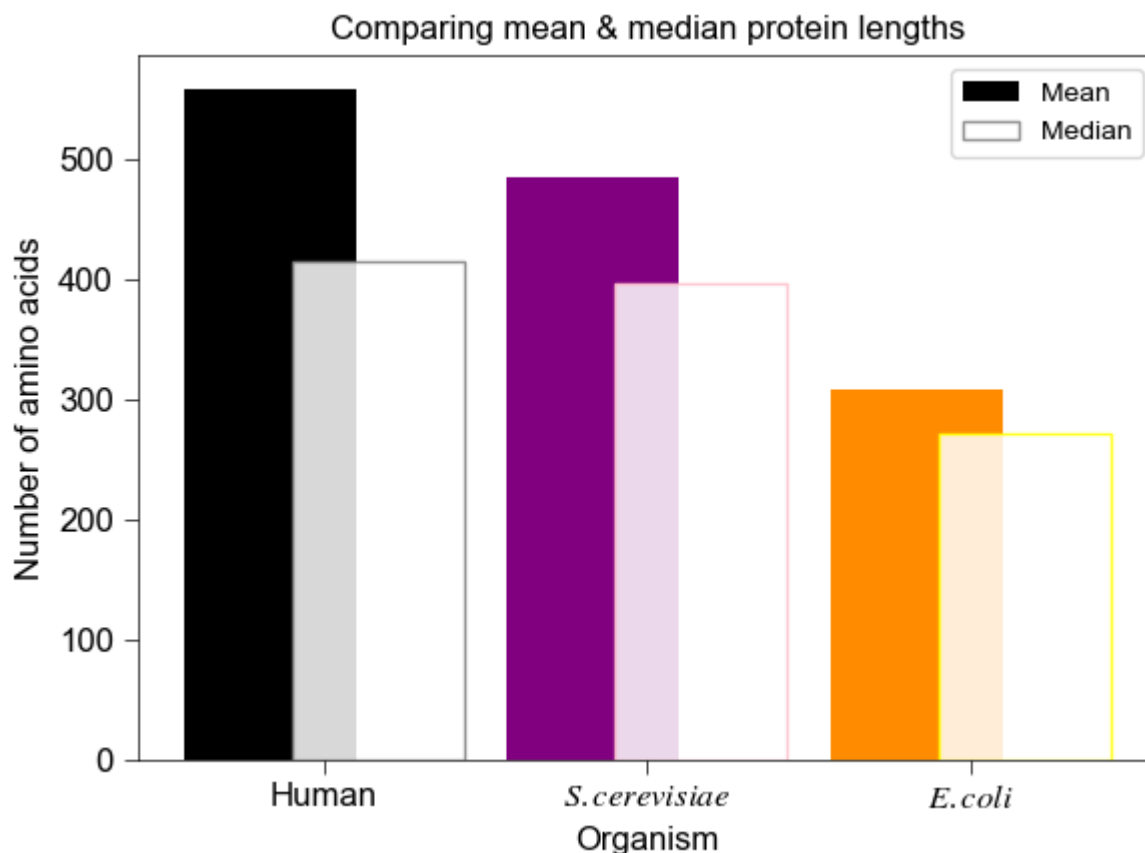
First, we see from the UniProt download pages that these two proteomes are much smaller than the human proteome: *E. coli* and *S. cerevisiae* respectively contain ~4,400 and ~6,000 proteins, which pale in number compared to the ~20,400 found in humans. Thus, to plot these distributions on the same y-axis, we must set the **density** option to **True** in our histograms, which ensures that the area under the histogram sums to 1. This enables us to compare histograms on the same y axis.

We can set up our script to compare these distributions:

```
1 # Instantiate the class
2 human = Proteome()
3
4 # Load the FASTA file containing the human proteome
5 uniprot = human.load_uniprot('uniprot_human.tab') # replace with the name of your .tab
  file
6 ecoLi = human.load_uniprot('uniprot_ecoli_k12.tab')
7 yeast = human.load_uniprot('uniprot_yeast.tab')
8
9 # How many amino acids are in each proteome?
10 print('\n')
11 print('Human has %.0f AA' % human.proteome_length(uniprot))
12 print('Yeast has %.0f AA' % human.proteome_length(yeast))
13 print('E coli has %.0f AA' % human.proteome_length(ecoLi))
14 print('\n')
15
16 # What are the mean & median lengths of a protein in each proteome?
17 print('Mean (median) human protein = %.0f AA (%.0f AA)' %
  (human.proteome_length(uniprot)/(len(uniprot['Length'][1:])),
  np.median(uniprot['Length'][1:]))
18 print('Mean (median) yeast protein = %.0f AA (%.0f AA)' %
  (human.proteome_length(yeast)/(len(yeast['Length'][1:])), np.median(yeast['Length']
  [1:]))
19 print('Mean (median) E coli protein = %.0f AA (%.0f AA)' %
  (human.proteome_length(ecoLi)/(len(ecoLi['Length'][1:])), np.median(ecoLi['Length']
  [1:]))
20
```

```
> Human has 11354232 AA
> Yeast has 2936363 AA
> E coli has 1354187 AA
> Mean (median) human protein = 558 AA (415 AA)
> Mean (median) yeast protein = 485 AA (396 AA)
> Mean (median) E coli protein = 309 AA (271 AA)
```

From this analysis, we can see a couple of interesting features. First, the average human protein (558 AA) is longer than the average yeast protein (485 AA) and nearly twice as long the *E. coli* average (309 AA). Interestingly, however, the median protein sizes from human and yeast are roughly equivalent, which suggests that the human proteome has a long(er) tail that includes some very large proteins.



Comparing the mean and median protein lengths across three different organisms

```
1 # Plot the mean & median protein lengths in a bar graph
2 fig = plt.figure(figsize=(6,4.5))
3 ax = fig.add_subplot(111)
4 ax.bar(1, human.proteome_length(uniprot)/(len(uniprot['Length'][1:])), color='k',
5 label='Mean')
6 ax.bar(1.5, np.median(uniprot['Length'][1:]), color='w', edgecolor='grey', alpha=0.85,
7 label='Median')
8 ax.bar(2.5, human.proteome_length(yeast)/(len(yeast['Length'][1:])), color='purple')
9 ax.bar(3, np.median(yeast['Length'][1:]), color='w', edgecolor='pink', alpha=0.85)
10 ax.bar(4, human.proteome_length(ecoli)/(len(ecoli['Length'][1:])), color='darkorange')
11 ax.bar(4.5, np.median(ecoli['Length'][1:]), color='w', edgecolor='yellow', alpha=0.85)
12
13 # Axis labels, legend, and title
14 ax.set_ylabel('Number of amino acids')
15 ax.set_xlabel('Organism')
16 ax.legend(loc='upper right')
17 ax.set_title('Comparing mean & median protein lengths')
```

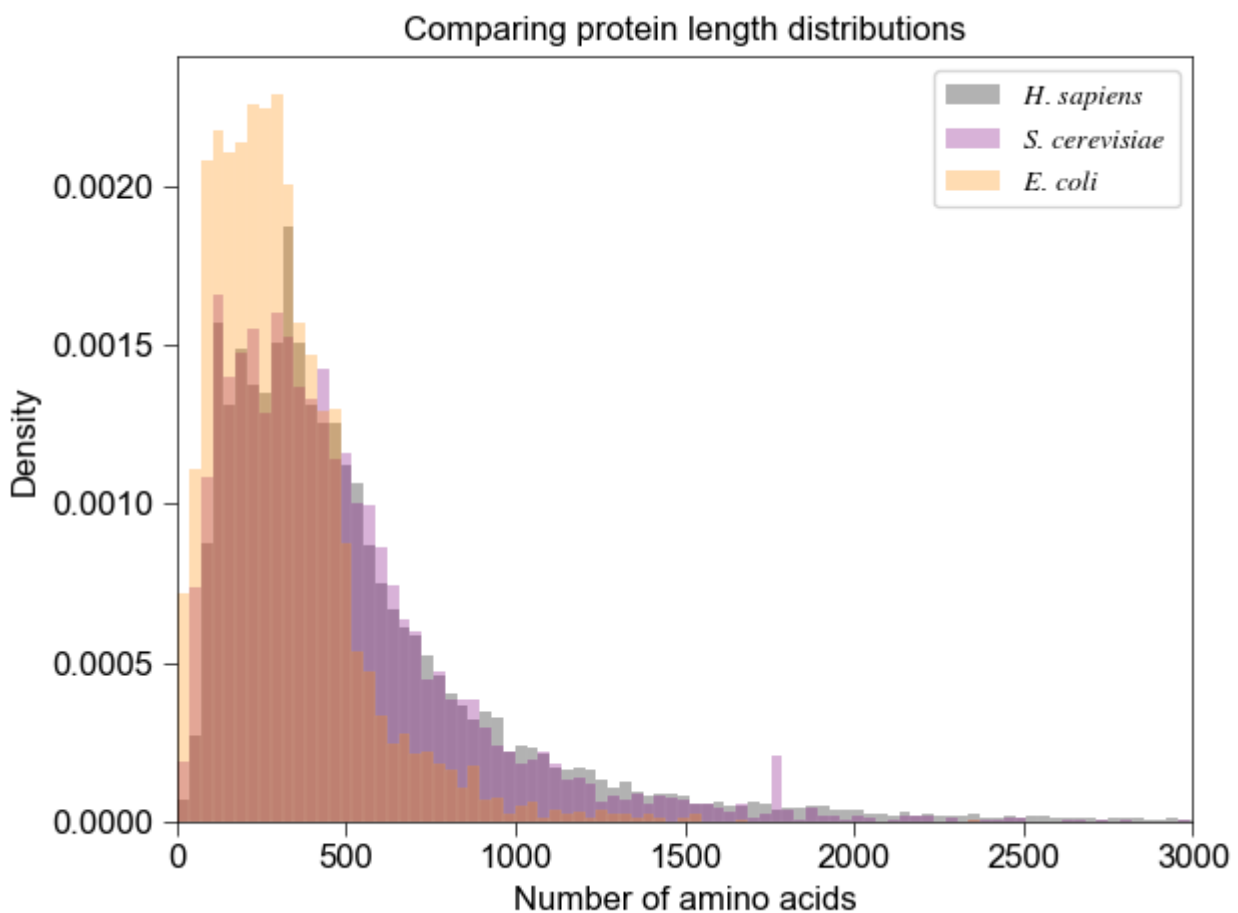


```

16
17 # Tick positions & labels
18 ax.set_xticks([1.25, 2.75, 4.25])
19 ax.set_xticklabels(['Human', r'$S$. $cerevisiae$', r'$E$. $coli$'])
20
21 plt.tight_layout()
22 plt.show()

```

We can also plot these data in histogram format as follows:



Comparing the distributions of protein lengths for three different organisms

```

1 # Get the bins from the human distribution
2 human_dat, human_bins = np.histogram(np.asarray(uniprot['Length'][1:]), bins=1000,
3 density=True)
4
5 # Plot the histograms using the same bins
6 plt.hist(uniprot['Length'][1:], bins=human_bins, color='k', alpha=0.3, edgecolor='w',
7 linewidth=0.1, label=r'$H$. $sapiens$', density=True)

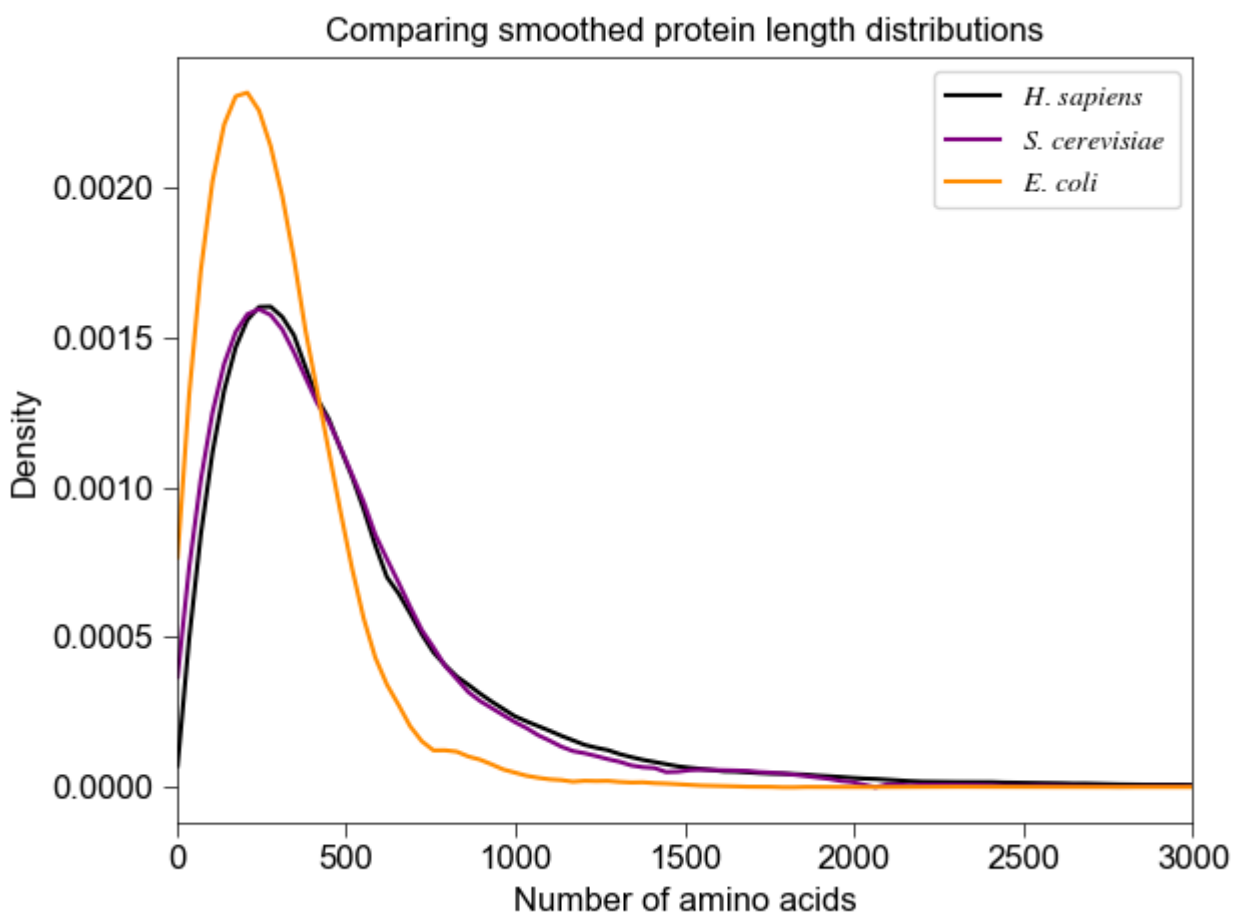
```

```

6 plt.hist(yeast['Length'][1:], bins=human_bins, color='purple', alpha=0.3,
edgecolor='w', linewidth=0.1, label=r'$S$. $cerevisiae$', density=True)
7 plt.hist(ecoli['Length'][1:], bins=human_bins, color='darkorange', alpha=0.3,
edgecolor='w', linewidth=0.1, label=r'$E$. $coli$', density=True)
8
9 # Set up the axis limits, title, and legend
10 plt.xlim(0,3000)
11 plt.ylabel('Density')
12 plt.xlabel('Number of amino acids')
13 plt.title('Comparing protein length distributions')
14 plt.legend(loc='upper right')
15
16 # plot
17 plt.tight_layout()
18 plt.show()

```

or with our smoothing function:



Comparing the smoothed distributions of protein lengths for three different organisms

```

1 # Histogram the data using the same bins
2 human_dat, human_bins = np.histogram(np.asarray(uniprot['Length'][1:]), bins=1000,
   density=True)
3 ecoli_dat, ecoli_bins = np.histogram(np.asarray(ecoli['Length'][1:]), bins=human_bins,
   density=True)
4 yeast_dat, yeast_bins = np.histogram(np.asarray(yeast['Length'][1:]), bins=human_bins,
   density=True)
5
6 # Plot the smoothed data as lines
7 plt.plot(human_bins[:-1], filter(human_dat), 'k-', label=r'$H$. $sapiens$')
8 plt.plot(yeast_bins[:-1], filter(yeast_dat), '-', color='purple', label=r'$S$.
   $cerevisiae$')
9 plt.plot(ecoli_bins[:-1], filter(ecoli_dat), '-', color='darkorange', label=r'$E$.
   $coli$')
10
11 # Set up the axis limits, title, and legend
12 plt.xlim(0,3000)
13 plt.ylabel('Density')
14 plt.xlabel('Number of amino acids')
15 plt.title('Comparing smoothed protein length distributions')
16 plt.legend(loc='upper right')
17
18 # Plot
19 plt.tight_layout()
20 plt.show()

```

For me, it was surprising to see how much the yeast proteome resembles the human proteome! Maybe to an evolutionary biologist this would not be surprising; however, to a structural biologist like myself, I never would have guessed that the proteome of the tiny, single-celled *S. cerevisiae* could look so similar to humans. Alongside protein length, perhaps the hubristic tendencies of humans have also been evolutionarily conserved...

The histograms above tell us **three things**:

1. there are many more small proteins in *E. coli*
2. proteins became significantly longer between the appearance of *E. coli* and that of *S. cerevisiae*
3. the distribution of protein lengths has not changed much between *S. cerevisiae* and humans.

Above we compared the protein length distribution of a prokaryote, *E. coli*, to two eukaryotes, the single-celled organism *S. cerevisiae* and far more complex *H. sapiens*. **What do protein length distributions from other prokaryotes or archaea look like? How do proteins in human mitochondria compare?**

I downloaded the proteomes (using the above outline) from the following organisms:

1. *Methanocaldococcus jannaschii* (an archaeon)
2. *Thermus thermophilus* (a bacterium)
3. *Bacillus subtilis* (a bacterium)

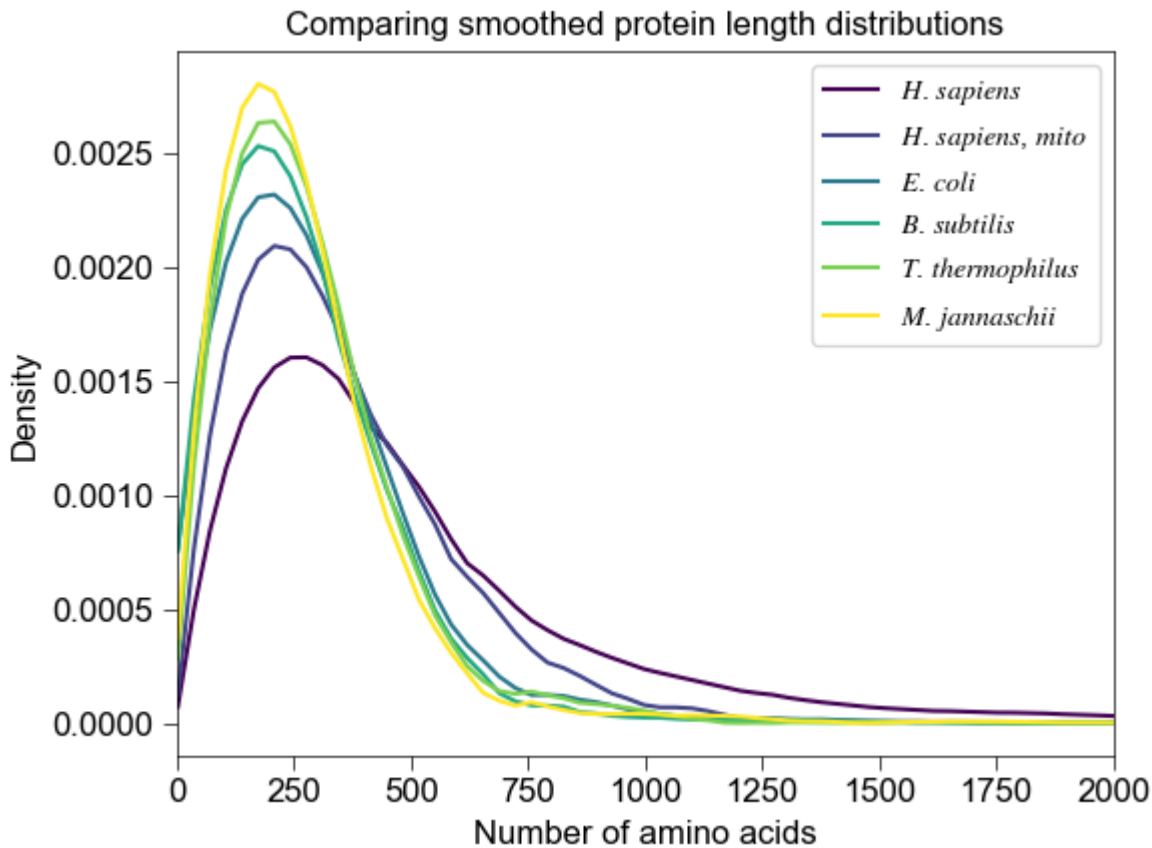
4. and the mitochondrial proteome from *Homo sapiens*

The mitochondrial proteome was made available by [The Broad Institute](#) in their MitoCarta2.0 release, which can be found in the following [publication](#).

```
1  # Load the new proteomes
2  mj = human.load_uniprot('m_jannaschii.tab')
3  bs = human.load_uniprot('b_subtilis.tab')
4  tt = human.load_uniprot('t_thermophilus.tab')
5
6  # in order to load the mitochondrial proteome, we need a new function
7  # that can read generic FASTA files and return a pandas dataframe
8
9  def load_fasta(self, fasta):
10     """ Reads a fasta file with identifiers (IDs) and sequences; returns a pandas
11     dataframe
12
13     Parameters
14     -----
15     fasta : file in FASTA format with identifiers (e.g. >) and sequences
16
17     Returns
18     -----
19     fasta_proteome, a Pandas dataframe that contains the FASTA IDs, sequences, and
20     length of sequences """
21
22     # Load the fasta file
23     self.seq_list = SeqIO.index(fasta, "fasta") # makes dictionary, but contains Seq
24     objects
25
26     # Create arrays to store the relevant information
27     ids = [] ; seqs = []
28     for key, value in self.seq_list.items():
29         ids.append(key)
30         seqs.append(str(value.seq)) # string the Seq object go get only the sequence
31
32     # Create a dictionary for subsequent conversion into a dataframe
33     self.dictionary = {'ID':ids, 'Sequence':seqs}
34
35     # Load the dictionary into a dataframe
36     self.fasta_proteome = pd.DataFrame(self.dictionary)
37     self.fasta_proteome['Length'] = self.fasta_proteome['Sequence'].str.len() #
38     create column in the dataframe with sequence lengths
39     self.fasta_proteome['Length'][1:] = self.fasta_proteome['Length']
40     [1:].astype(int) # convert the length values to integers
41
42     return self.fasta_proteome
43
44 # now load the mitochondrial proteome using "load_fasta"
45 mito = human.load_fasta('Human.MitoCarta2.0.fasta')
```

Now that the data have been loaded into pandas dataframes, we can plot the smoothed protein length distributions. To speed things up, we can write a function:

```
1 def plot_list_hist(list_organisms, names):
2
3     # get reference data
4     ref_dat, ref_bins = np.histogram(np.asarray(list_organisms[0]['Length'][1:]),
5     bins=1000, density=True) # get reference bins for 1st item in list
6
7     # make array
8     data = np.zeros(shape=(len(list_organisms), len(ref_bins)-1))
9     colors = plt.cm.viridis(np.linspace(0.0, 1, len(list_organisms)))
10
11     # loop and plot
12     fig = plt.figure(figsize=(6,4.5))
13     for i, item in enumerate(list_organisms):
14         dat, bins = np.histogram(np.asarray(item['Length'][1:]), bins=ref_bins,
15         density=True)
16         data[i, :] = dat
17         plt.plot(ref_bins[:-1], filter(dat), '-', label=names[i], color=colors[i])
18
19     plt.xlim(0,2000)
20     plt.ylabel('Density')
21     plt.xlabel('Number of amino acids')
22     plt.title('Comparing smoothed protein length distributions')
23     plt.legend(loc='upper right')
24     plt.tight_layout()
25
26     #plt.show()
27     return fig
28
29 # create a list of data and data names
30 data = [uniprot, mito, ecoli, bs, tt, mj]
31 data_names = [r'$H$. $sapiens$', r'$H$. $sapiens$, $mito$', r'$E$. $coli$', r'$B$.
32 $subtilis$', r'$T$. $thermophilus$', r'$M$. $jannaschii$']
33
34 # use the "plot_list" function
35 plot_list_hist(data, data_names)
36 plt.show()
```



Comparing the smoothed distributions of protein lengths across organisms

From this comparison, we can see that the **mitochondrial** proteome lies somewhere between the ***E. coli*** and **human** proteomes. Of course, our human proteome dataset also contains all of the mitochondrial proteins; so, had we selectively removed the ~1000 mitochondrial proteins from the human proteome, we would see a slightly more pronounced difference. Nevertheless, for proteins with fewer than 500 amino acids, the mitochondrial proteome looks more similar in terms of its length distribution to *E. coli*. For proteins longer than 500 residues, however, the long tail of the mitochondrial proteome bears more resemblance to human. Overall, these apparent differences may reflect the more ancient evolutionary origin of mitochondria: more smaller proteins than expected for the human proteome, but more larger proteins than expected for prokaryotes or archaea.

Finally, we can also plot the mean and median protein lengths, as we did earlier. To make this a little faster, however, we can create a small loop:

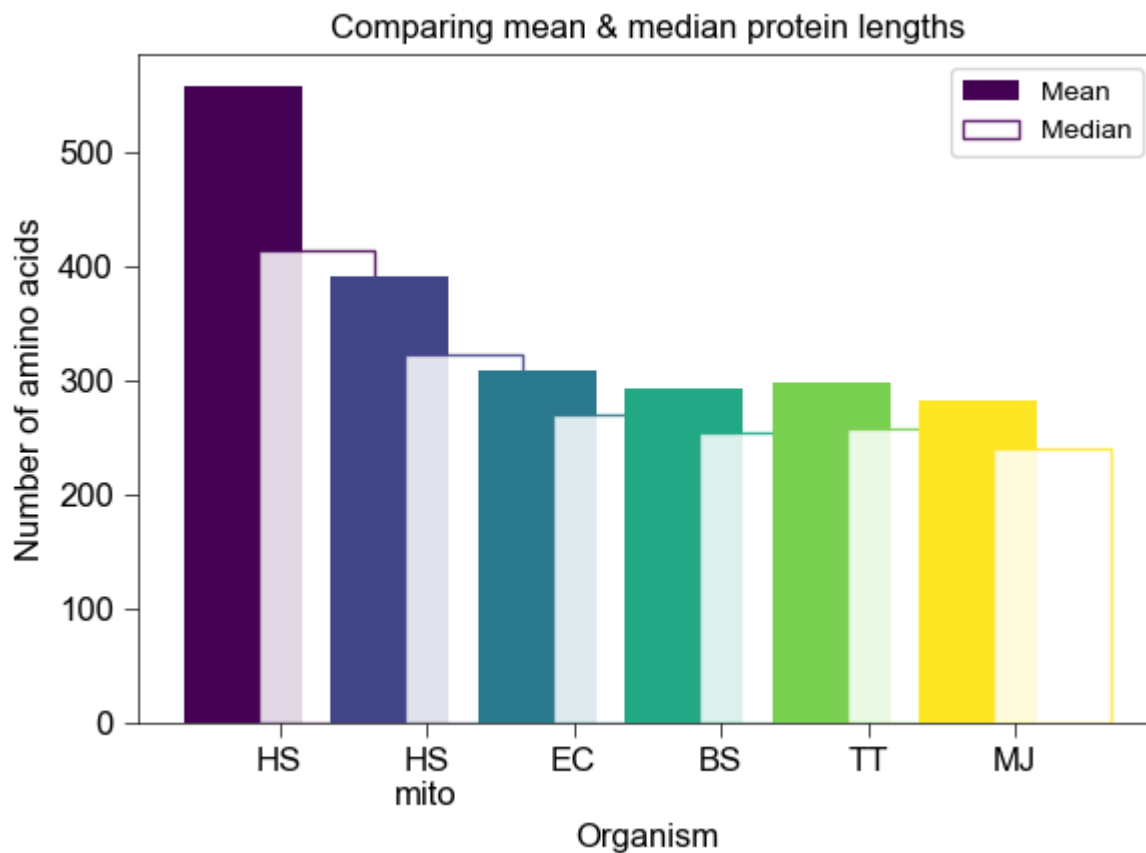
```
1 # create the figure
2 fig = plt.figure(figsize=(6,4.5))
3 ax = fig.add_subplot(111)
4
5 # create a list of colors using the colormaps package
6 abbrev_names = ['HS', 'HS\nmito', 'EC', 'BS', 'TT', 'MJ']
7 colors = plt.cm.viridis(np.linspace(0.0, 1, len(abbrev_names)))
8
9 # initiate a counter and empty array (for xtick positions)
```



```

10 counter = 1
11 xtick_list = []
12
13 # now loop over the organisms
14 for i in range(len(abbrev_names)):
15     if i == 0:
16         ax.bar(counter, human.proteome_length(orgs[i])/(len(orgs[i]['Length'][1:])),
17 color=colors[i], label='Mean')
18         ax.bar(counter+0.5, np.median(orgs[i]['Length'][1:]), color='w',
19 edgcolor=colors[i], alpha=0.85, label='Median')
20     else:
21         ax.bar(counter, human.proteome_length(orgs[i])/(len(orgs[i]['Length'][1:])),
22 color=colors[i])
23         ax.bar(counter+0.5, np.median(orgs[i]['Length'][1:]), color='w',
24 edgcolor=colors[i], alpha=0.85)
25         xtick_list.append(counter+0.25)
26         counter += 1
27
28 # Axis labels & the legend
29 ax.set_ylabel('Number of amino acids')
30 ax.set_xlabel('Organism')
31 ax.legend(loc='upper right')
32
33 # Tick positions & labels
34 ax.set_xticks(xtick_list)
35 ax.set_xticklabels(abbrev_names)
36 ax.set_title('Comparing mean & median protein lengths')
37
38 #plt.xticks(rotation=0, fontsize=8)
39 plt.tight_layout()
40 plt.show()

```



Comparing the mean and median protein lengths across organisms. The abbreviations are as follows: Homo sapiens (HS), mitochondrial (mito), E. coli (EC), Bacillus subtilis (BS), Thermus thermophilus (TT), Methanocaldococcus jannaschii (MJ).

We can clearly see that the median protein length in the mitochondria is ~100 residues smaller than the typical human protein. The difference in the means is even more pronounced, which reflects the very long tail in the human proteome (as discussed above).

Next time, I'll discuss how we can use Python to calculate biophysical properties from protein sequences, and how we can perform such calculations on the proteome scale.